

# Homework 2: PyTorch and Self-Attention

INFO 159 / 259, Spring 2023

**Due:** February 7, 2023 (11:59 PM)

---

This homework is designed to help you:

1. familiarize yourself PyTorch and implement a feedforward neural network (FFNN) [deliverable 1]
2. understand the self-attention mechanism that is crucial to the Transformer architecture [deliverable 2]

This document contains the instructions for this homework, and the accompanying notebook can be found here:

<https://github.com/dbamman/nlp23/blob/main/HW2/HW2.ipynb>

To use it, like in HW1, hit the “Open in Colab” button  at the top, and be sure to save a copy in your Google Drive or Github *before* you start to work on the file.

Submit your work on Gradescope. For code questions (Q2, Q4–5), replace `raise NotImplementedError` with your solutions, and as in HW1, keep your code between the `# BEGIN SOLUTION` and `# END SOLUTION` flags. For short-answer questions (Q1 and Q3), submit all your answers as a `.pdf` file separately (see §4 for instructions on how to submit).

## 1 Deliverable 1: PyTorch and FFNN

The first deliverable can be seen as a tutorial of PyTorch, which is a standard library that researchers use to implement neural networks.

**QUESTION 1** [10 pts] **PyTorch and embeddings** (writeup only). In the notebook, you can see that a PyTorch implementation of logistic regression is offered for your reference. Study this script carefully. Start with how we load and access the GloVe embeddings. The most critical ingredients of a neural net in PyTorch include a class that defines the architecture (pay attention to the `forward` method), an optimizer (`torch.optim.Adam`), and a loss function, `torch.nn.CrossEntropyLoss()`, which combines the softmax function `torch.nn.LogSoftmax()` and negative log-likelihood `torch.nn.NLLLoss()`.<sup>1</sup>

A notable difference from HW1 here is that the input of this model is the average of GloVe embeddings for all words in a movie review.<sup>2</sup> What is the difference between averaged embeddings and BoW? What are the advantages or disadvantages of each? Discuss in no more than 100 words.

**QUESTION 2** [10 pts] **Feedforward neural network (FFNN)**. For this question, we want to add a hidden layer to the logistic regression classifier you’ve seen in Q1. Implement Eqn. (7.13) in J&M SLP3 and let the non-linearity  $g$  be `tanh`. Your implementation should be similar to the `LogisticRegressionClassifier` example in the notebook. Let’s pick 20 for the size of the hidden layer—this is provided in the `__init__()` function

---

<sup>1</sup>See documentation: <https://pytorch.org/docs/stable/generated/torch.nn.CrossEntropyLoss.html>

<sup>2</sup>For more on GloVe embeddings, see: <https://nlp.stanford.edu/projects/glove/>

(hidden\_dim=20), so you don't need to worry about it. Note that in the J&M terminology, a “two-layer” network has one hidden layer, which is what you will be implementing. You should fill in the parts between “#BEGIN SOLUTION” and “#END SOLUTION.”

## 2 Deliverable 2: Self-attention basics

The self-attention mechanism is often thought of as one of the most transformative ideas in modern NLP. Its full form in Transformer, as introduced in “Attention is All You Need,”<sup>3</sup> is rather involved. This deliverable aims to prepare you for it.

We will start with the simplest form of self-attention: scaled dot-product self-attention. The goal is to try to understand the roles that query, key, and value vectors play in attending to the input sequence: conceptually, what do they aim to achieve and improve on? How do you code this in Python? Instead of focusing on training models, here we really want to see vector transformations in action.

**QUESTION 3** [10 pts] **The concept of self-attention** (writeup only). Self-attention may be one of those less intuitive concepts you have encountered in this course. In the era of ChatGPT, you might be curious if the chatbot can *tutor* you and answer your questions about a difficult concept in NLP. This question is your chance to play around with it and start thinking about the power and limitations of such tools. First, make sure that you have worked through §10.1 in SLP 3 and the relevant lecture on self-attention. Then, you can go to <https://chat.openai.com/chat>, ask ChatGPT questions, and most importantly, evaluate its answers. Here's an example prompt:

Why do we need the self-attention mechanism in NLP?

It yields the following answer:

Self-attention is a mechanism in deep learning models for processing sequences, where each element in a sequence attends to all other elements in a weight-based manner to compute its representation. This enables the model to dynamically weigh the importance of different elements in the sequence for the current task, and allows for more context-aware representations compared to traditional recurrent neural networks.

Is this correct? Substantiate your reasoning in less than 100 words.

**QUESTION 4** [10 pts] **Scaled dot-product attention in NumPy.** In the notebook, we'll go through the basics of implementing the steps in attention outside of any model. We'll do that in NumPy to help you see the changes in dimensions over this process. Then, for this question, implement the attention function below in NumPy. Recall from lecture that attention given a query vector  $Q$ , key vector  $K$  and value vector  $V$  is given by the following equation:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V$$

Where  $Q$ ,  $K$  and  $V$  are the query, key and value matrices created by multiplying an input embedding matrix (for a sentence) by  $W^Q$ ,  $W^K$  and  $W^V$ , respectively. You will calculate  $Q$ ,  $K$ ,  $V$  within the body of this function. The sole required argument to this function should be a 2-D matrix  $\in \mathbb{R}^{n \times \text{input\_embedding\_size}}$  for any arbitrary  $n$  (that is, corresponding to a sentence of arbitrary length). It should return a matrix of that same exact size that is the output of that attention process over the input, given the parameters specified below.  $d_k$  here is the size of the key vector (query\_key\_size=37).

*Hints:*

- When taking softmax, remember to operate on the correct dimension. Work through the “Intuition of attention” section if you're unsure.

---

<sup>3</sup>Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, Illia Polosukhin, “Attention Is All You Need” (NIPS 2017).

- Use `a.shape` or `a.size()` to view the dimensions of a torch variable `a`, which you'll need to get  $d_k$ .
- For matrix multiplication between `A` and `B`, use `A @ B`.

**QUESTION 5** [10 pts] **Scaled dot-product attention in PyTorch.** Now it's time to implement that as part of a model. Here we're going to embed attention within a larger model. For an input document of, say, 20 words, each represented by a 100-dimensional embedding, the input to attention is a  $20 \times 100$  matrix; the output from attention is also a  $20 \times 100$  matrix. In this larger model, we're going to average those output embeddings to generate a final document that's a single 100-dimensional vector; pass through a fully-connected dense layer to make a prediction.

You will finish the implementation of the `Attention` class in the notebook. When used in a model (as distinct from the NumPy example above), it is important to note that  $W^Q$ ,  $W^K$  and  $W^V$  are all trainable (this is how the attention mechanism learns). You'll see in the provided code that the output of your attention mechanism is averaged over all token embeddings and then fed into a learnable linear layer to make a categorical prediction—that's all taken care of for you, so you only need to implement the attention module. Since  $W^Q$ ,  $W^K$  and  $W^V$  are matrices, initialize them as a linear layer (`nn.Linear`) without the learnable additive bias (i.e., `bias=False`). You can also initialize the softmax layer in the `__init__()`. The `forward()` method should be a very straightforward implementation of the formula above and can be done within ten lines of code.

### 3 Debugging guides

For this homework, you might find it helpful to use these tips to help debug your code:

- The PyTorch function `.shape` can be called on torch variables to view their dimensions. This might be helpful in ensuring the layer dimensions are set correctly.
- Python debugger is a valuable tool which you can use to step through your program's execution. The line of code `"import pdb; pdb.set_trace()"` will add a breakpoint to your code, where you can view dimensions of PyTorch layers and experiment with function calls.

### 4 How to submit

Submit your work to Gradescope. There are two different Gradescope submission links, one for the PDF writeup and one for code:

- Submit to: "Gradescope HW2 writeup"
  - Submit all your answers to Q1 and Q3 as one PDF file to Gradescope. As long as you submit a PDF, there's no additional requirements for formatting. There is no template for writeups.
- Submit to: "Gradescope HW2 code (ipynb)"
  - Make sure you have replaced all `raise NotImplementedError` with your solutions, and that they are placed, as in HW1, between the `# BEGIN SOLUTION` and `# END SOLUTION` flags.
  - Download your Colab notebook as an `.ipynb` file (File → Download `.ipynb`)
  - Submit `HW_2.ipynb`. The file must be named in this way for the Gradescope autograder.